


PruneVid question-to-visual attention across LLM layers (PLLaVA-7B). Tokens are pruned after layer 10 ($\alpha=0.4$), so layers 20/30 attend only to kept tokens.


frames



Layer 5



Layer 10



Layer 20



Layer 30



Q: What is this video about? A: The video appears to be a scene from a movie or television show featuring a character who is speaking in a language that uses Chinese charac