


PruneVid question-to-visual attention across LLM layers (PLLaVA-7B). Tokens are pruned after layer 10 ( $\alpha=0.4$ ), so layers 20/30 attend only to kept tokens.

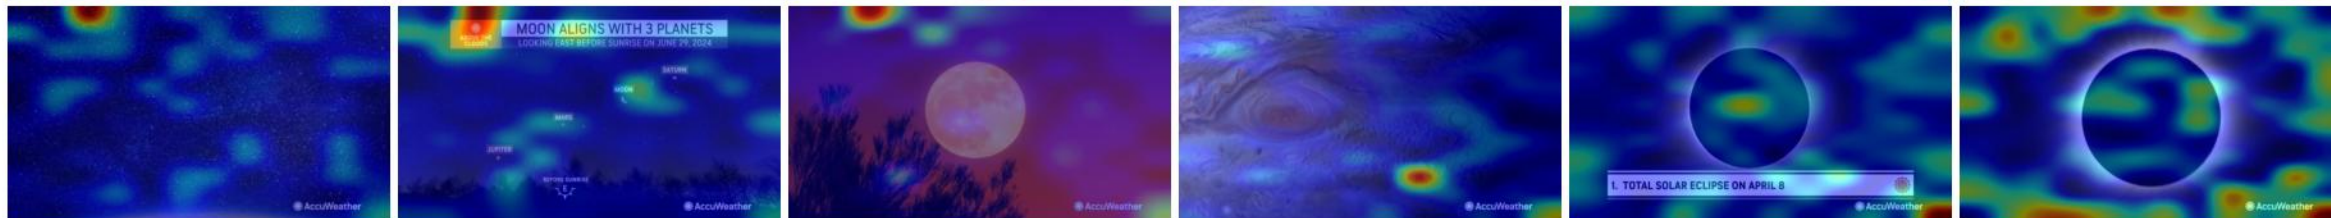
frames



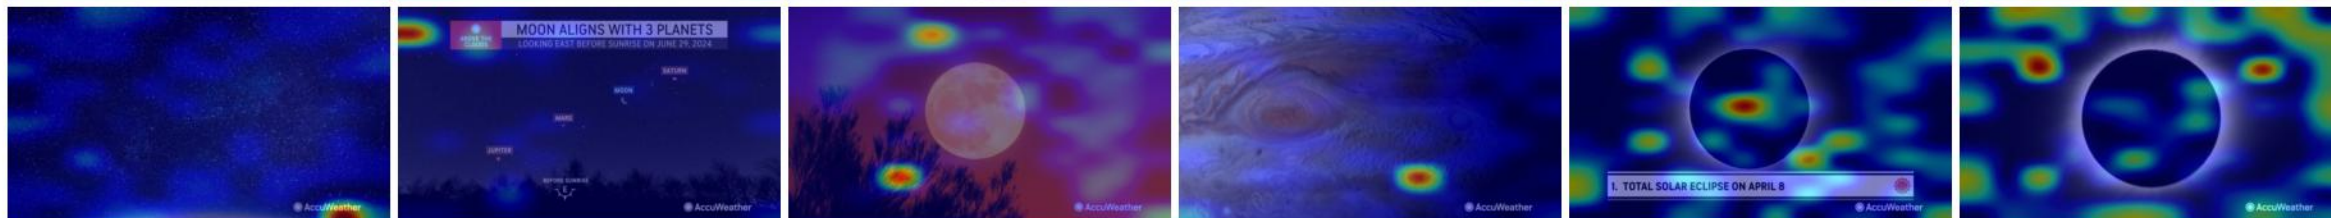
Layer 5



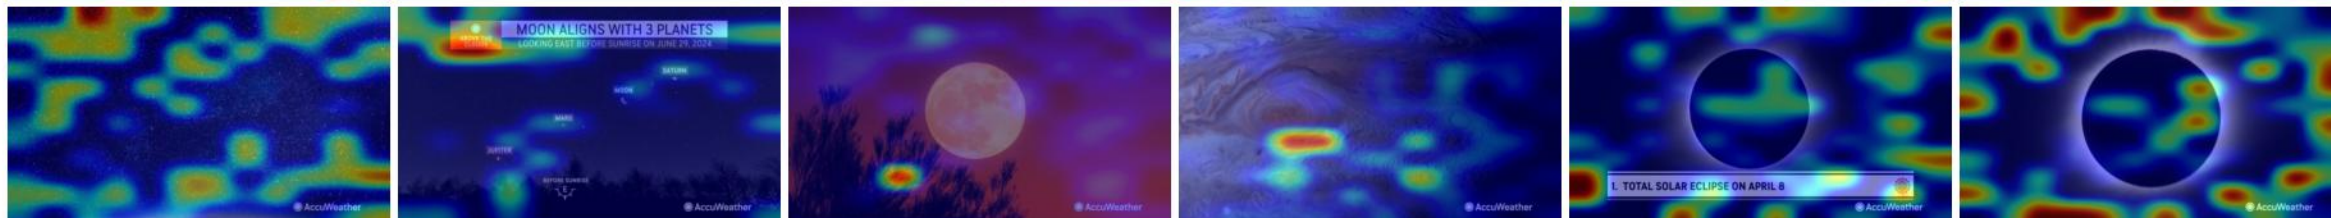
Layer 10



Layer 20



Layer 30



Q: What is this video about? A: The video appears to be a simulation or animation showing the path of the total solar eclipse on April 8, 2024. The eclipse is depicted mov