

PruneVid (PLLaVA-7B) | blue=static, orange=dynamic (tau=0.8) | layer=10 question-to-visual attention | kept 438/1110 merged tokens (segments: [5, 3, 6, 2])


frames



static/dyn



attention



kept



Q: What is this video about? A: The video appears to be a scene from a movie or television show featuring a character who is speaking in a language that uses Chinese charac